

Computer Science Department

CS660 – Mathematical Foundations of Analytics (CRN# 71425)

Fall 2024

Project #3 / Due 09-Dec-2024

Let's review how the **CART** (Classification And Regression Tree) algorithms work in Python's scikit-learn module.

Namely, we should study the following **Decision Tree** libraries:

1_ DecisionTreeClassifier

<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

2_ DecisionTreeRegressor

<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeRegressor.html>

3_ RandomForestClassifier

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>

4_ RandomForestRegressor

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>

You should perform both Classification and Regression modeling on the following datasets:

A_ For **classification** use scikit-learn's hand-written digits (each data point is an 8X8 image of a single digit)

https://scikit-learn.org/stable/modules/generated/sklearn.datasets.load_digits.html#sklearn.datasets.load_digits

```
>>> from sklearn.datasets import load_digits
>>> digits = load_digits()
>>> print(digits.data.shape)
(1797, 64)
```

B_ For **regression** use scikit-learn's California-housing dataset

https://scikit-learn.org/stable/modules/generated/sklearn.datasets.fetch_california_housing.html#sklearn.datasets.fetch_california_housing

```
>>> from sklearn.datasets import fetch_california_housing
>>> df_cal_housing = fetch_california_housing(as_frame=True)
```

The target feature is 'MedHouseVal'. This is the value of a house; you need to predict it.

```
>>> df_cal_housing.frame.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MedInc          20640 non-null  float64
1   HouseAge        20640 non-null  float64
2   AveRooms        20640 non-null  float64
3   AveBedrms       20640 non-null  float64
4   Population      20640 non-null  float64
5   AveOccup        20640 non-null  float64
6   Latitude        20640 non-null  float64
7   Longitude       20640 non-null  float64
8   MedHouseVal     20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
>>> df_cal_housing.frame.head()
   MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  Longitude  MedHouseVal
0   8.3252     41.0   6.984127   1.023810     322.0   2.555556     37.88    -122.23     4.526
1   8.3014     21.0   6.238137   0.971880    2401.0   2.109842     37.86    -122.22     3.585
2   7.2574     52.0   8.288136   1.073446     496.0   2.802260     37.85    -122.24     3.521
3   5.6431     52.0   5.817352   1.073059     558.0   2.547945     37.85    -122.25     3.413
4   3.8462     52.0   6.281853   1.081081     565.0   2.181467     37.85    -122.25     3.422
```

Write **Python/R** scripts (within a Notebook) in order to complete the following tasks:

- 1_ Perform Exploratory Data Analysis (EDA) on both datasets.
- 2_ Perform classification task on the digit's dataset, utilizing both DTClassifier and RandomForest with their default parameters. Which one performs better? Print out their respective Confusion Matrices.
- 3_ Perform same task as #1, but now tune the classifiers. Which one performs better? Print out their respective Confusion Matrices.
- 4_ Perform regression analysis task on the California-housing dataset, utilizing both DTRRegressor and RandomRegressor with their default parameters. Which one performs better? Print out their respective Mean-Squared Error (MSE) and Coefficient of Determination (R-squared, R^2)
- 5_ Perform same task as #3, but now tune the regressors. Which one performs better? Print out their respective Mean-Squared Error (MSE) and Coefficient of Determination (R-squared, R^2)
- 6_ Print out, for each of the tasks above, the feature importance (aka Gini importance) for each of the features in the dataset.

property feature_importances_

Return the feature importances.

The importance of a feature is computed as the (normalized) total reduction of the criterion brought by that feature